

Wearables and Habits: Data Audit, Variable Construction and Balance

Issue #5

4 October 2026

1 Summary

Exploratory audit of Open e-commerce 1.0. Main takeaways:

- **Usable window is 2018-01 to 2023-03, not 2024-08.** Users are observed from their first to last purchase, and almost every window ends by early 2023 (a single user remains after 2023-03). Calendar series are restricted to months with $\geq 1,000$ active users.
- **Treatment:** 1,016 users buy a tracker or smartwatch (Amazon categories `WEARABLE_COMPUTER`, `BIOMETRIC_MONITOR`); 621 have ≥ 12 months of pre and post history, 179 have ≥ 24 .
- **Series are smooth in levels** after the restriction (no month-to-month jumps beyond 3 sd except for one running-shoe month), but rare-event counts (wearable, running shoes, fitness gear) are noisy month to month, as expected from small counts.
- **Balance fails.** Ever-treated users differ from never-treated users on income, education, household size and, strongly, on how long they are observed and how much they buy. Adopters are also on a rising purchasing trend before adoption. This supports the plan's main threat (selection and anticipation); the design must use never-treated/not-yet-treated comparisons conditional on covariates, and event-study pre-trends are the key test.
- The running proxy is only a purchase proxy and is weak in count terms (about 1 in 100 user-months has a running-shoe purchase).

2 Audit

	Value
rows	1,850,717
users	5,027
survey users	5,027
purchase users not in survey	0
missing Category (%)	4.83
missing Title (%)	4.85
missing price (%)	0.00
missing qty (%)	0.00
price ≤ 0 rows	0
qty > 1 rows (%)	5.57
exact duplicate rows	11,624
date min	2018-01-01
date max	2024-08-15

Quirks worth carrying forward: 4.8% of rows have no Category or Title (they cannot be classified and are treated as non-flagged); 11,624 rows are exact duplicates (could be genuine repeat orders, kept, not deduplicated); 5.6% of rows have quantity above one (spend uses price $\times$ quantity). All 5,027 purchasers are in the survey.

3 Variables built

The user-month panel (`code/build/interm/panel_user_month.pkl`, gitignored) spans each user's first to last purchase month with zeros filled; 273,483 user-months.

- **Treatment:** first month of a purchase in `WEARABLE_COMPUTER` or `BIOMETRIC_MONITOR`. Watch bands and plain `WATCH` are excluded (the latter mixes in non-smart watches). Kids' trackers are not yet excluded.
- **Running-shoe:** category `SHOES` or `TECHNICAL_SPORT_SHOE` with a running keyword in the title (2,580 rows).
- **Running-related** (broader): any running keyword in the title outside table runners, books and movies, plus `TECHNICAL_SPORT_SHOE` (13,456 rows). It includes socks, shorts and head-phones, so it is loose; keep both versions.
- **Fitness gear, supplements** (secondary or placebo-type outcomes), total purchases and spend (activity controls). Each is a count and a > 0 dummy per user-month.

	Mean	Sd	Pct zero		
				Users	
Purchases	6.767	10.747	27.035	All users	5027
Spend (\\$)	161.084	292.261	27.035	Ever bought wearable	1016
Wearable purchases	0.006	0.086	99.464	$\geq 6m$ pre-history	927
Running-related purchases	0.049	0.283	96.123	$\geq 12m$ pre and post	621
Running-shoe purchases	0.009	0.118	99.209	$\geq 24m$ pre and post	179
Fitness-gear purchases	0.025	0.189	97.893		
Supplement purchases	0.138	0.559	91.162		

The wearable share is tiny per user-month, so any event-study on a single month will be noisy; bins of 3 months or a dummy outcome over a longer window will probably be needed.

4 Aggregate series and smoothness

Series	Mean	Sd(diff)/Sd(level)	Months $ z > 3$	AR(1)
Purchases per active user	6.71	0.63	0	0.81
Spend per active user (USD)	159.83	0.72	0	0.74
Wearable purchases per 1000 users	5.77	1.17	0	0.32
Running-related per 1000 users	48.42	0.62	0	0.81
Running-shoe per 1000 users	9.38	0.99	1	0.51
Fitness-gear per 1000 users	24.97	0.93	0	0.57
Supplement per 1000 users	138.42	0.51	0	0.87
Active users	4,340.89	0.29	2	0.95

All series trend up from 2018 to 2021–22, consistent with growth in purchasing within panel members (and the composition of the panel). Rates are smooth except for the count-based wearable, running-shoe and fitness-gear series, where $sd(\Delta)/sd(level)$ is close to one; that is sampling noise,

not data breaks. Before restricting to $\geq 1,000$ active users, the tail had spikes and zeros that come entirely from one or two remaining users; that tail is excluded. The active-user count itself drops sharply after 2022-11, so months after that have a rapidly changing sample.

Event-time series show (i) a mechanical spike at $t = 0$ in total purchases, spend and running-related purchases, because the purchase of the device itself is counted and baskets are bigger around that month, and (ii) a rising pre-adoption trend in purchases, spend, running-related and supplement purchases. The $t = 0$ month should be dropped or modeled separately, and the pre-trend must be handled by the control group and the estimator (not-yet-treated comparison). Running-shoe and fitness-gear series are too noisy to read at monthly frequency.

5 Correlations

Everything is positively correlated with overall purchasing intensity (running-related 0.61, supplements 0.62, fitness gear 0.49 with total purchases). The wearable–running correlation (0.28) is no larger than wearable’s correlation with supplements (0.23) or fitness gear (0.24), so there is no evident running-specific link at the user level; the raw association is mostly “heavy buyer” selection, and outcomes should be conditioned on or scaled by overall activity.

6 Balance tests

Linear probability models (HC1 standard errors, covariates standardized so coefficients are the change in $P(\text{ever treated})$ per +1 sd) with joint F -tests. Demographics come from the survey; income “prefer not to say” and other missing answers drop out, so $N = 3,761$ for the demographic regressions.

Regression	Joint F p-value	N
Ever-treated on demographics	0.000	3,761
Ever-treated on demographics + log months	0.000	3,761
Ever-treated on first-12m purchasing	0.000	4,968
Adoption date on demographics (treated only)	0.006	811

All joint tests reject. Treated users are higher income, more educated or older, in larger households, and observed for longer (a +1 sd longer window is worth about +4 percentage points of adoption probability against a base of about 20%, which is partly mechanical: more months give more chances to buy). They also buy more overall in their first year. Conditional on overall purchasing and spend, early running, fitness and supplement rates do not predict adoption, but unconditionally ever-treated users buy about twice as many running-related items as never-treated users throughout the sample (the running-related figure above), so the level difference is largely general buying intensity. Cohorts are also not balanced on demographics ($p = 0.006$), so timing is not as-good-as-random.

7 Go/no-go read and next steps

- Sample size is workable for a reduced-form design (about 600 treated users with ≥ 12 months pre and post) but thin for monthly running-shoe outcomes. Preliminary read: **conditional go**, with the broader running measure and quarterly bins as the primary outcome.
- Needed next: exclude kids’ trackers and gift/duplicate-looking purchases; decide on the running outcome definition (tighter keyword audit of titles); set up the Callaway–Sant’Anna

comparison group (not-yet-treated plus never-treated) with demographics and pre-adoption purchase intensity as covariates; handle the $t = 0$ spike and the 2022-11 sample drop.

- Environment note: the compose `container_name: research-container` collides with another project's running container on this machine (mounted from `exploration-exploitation`), so this audit was run with `docker compose run` on this repo's `code-opt` image (same service, no new container). R `fixest` and `did` load there; the audit regressions are in Python `statsmodels` because they are simple LPMs.




Figure 1: Calendar-month series per active user (or per 1,000 active users) for each constructed variable, with centered 12-month moving average.




Figure 2: Ever-treated users: mean outcomes by months since first wearable purchase (± 24 months; composition changes across event times).




Figure 3: Running-related purchases per 1,000 user-months, ever-treated versus never-treated (6-month moving average).




Figure 4: Spearman correlations of per-user monthly rates (full observed window).




Figure 5: Ever-treated vs never-treated on survey demographics, with and without log months observed.




Figure 6: Ever-treated on purchasing rates in each user's first 12 observed months.




Figure 7: Among ever-treated users: adoption date on demographics (balance across cohorts).